

# Meta Learning for System Identification

Marco Forgione

In collaboration with Dario Piga, Matteo Rufolo, Riccardo Busetto, Gabriele Maroni, Filippo Pura,...

## Motivation


- Machine learning models can effectively describe a wide variety of dynamical systems
- However, training is computationally expensive and data-hungry
- Can we **automate the design of algorithms**, model structures, optimizers **for a class** of system identification problems?
- This is meta learning, aka learning to learn!

## The “sines” meta learning toy example

- Regression datasets  $x \rightarrow y$ : sines with randomized phase and amplitude

$$y(x) = A \sin(x + \psi)$$

- Simple structure, but just  $K = 5$  points per dataset
- Would be trivial with  $K = 50$  points per dataset, or if we the knew the true model structure

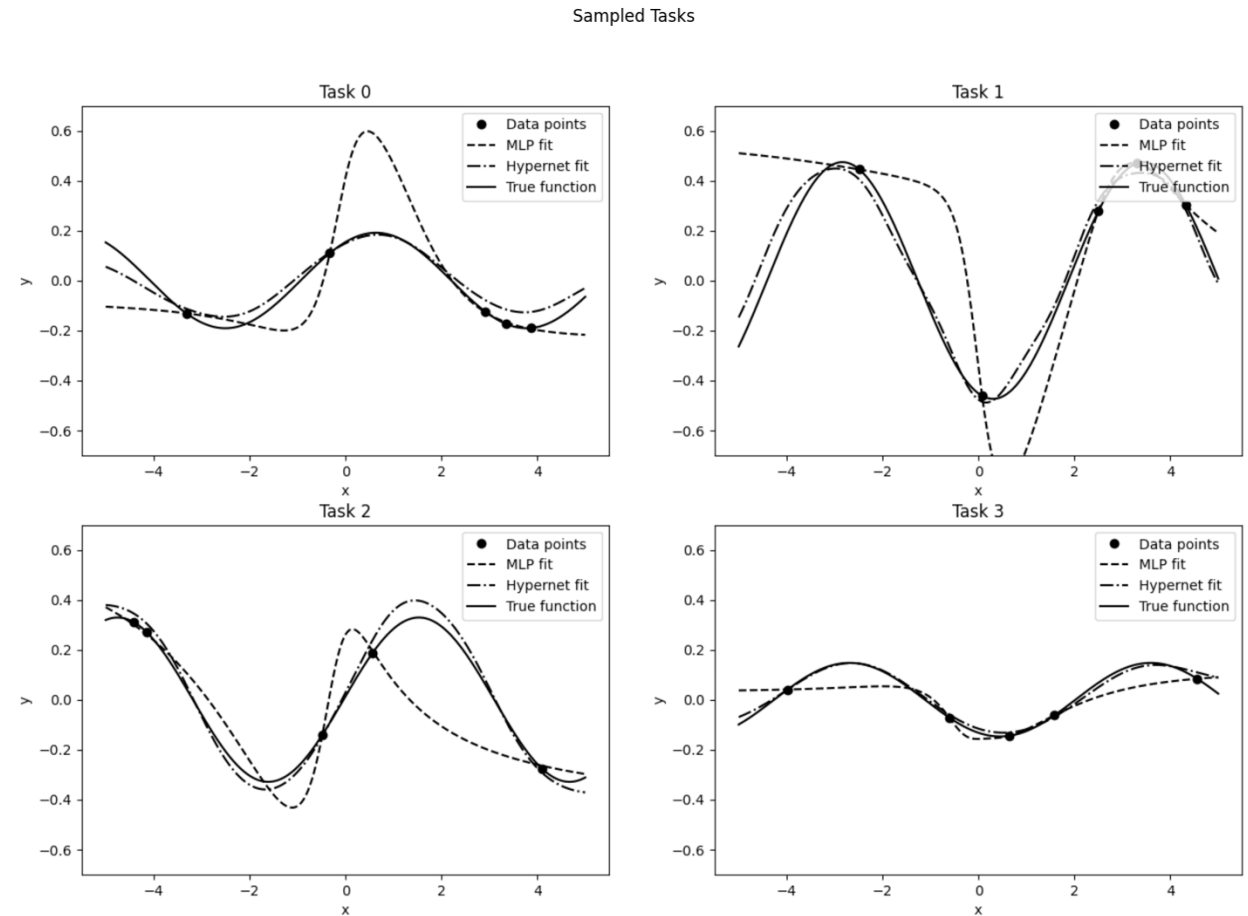


Example from:

Finn, Chelsea, Pieter Abbeel, and Sergey Levine. "Model-agnostic meta-learning for fast adaptation of deep networks." International conference on machine learning. PMLR, 2017.

## The “sines” meta learning toy example

- With a black-box structure (eg, a small neural net) we overfit the  $K = 5$  data points
- If we knew the structure,  $K = 2$  data points would be enough to find  $A, \psi$
- What if we observe  $K = 5$  points from many datasets? Can we learn the pattern?
- Results with existing meta learning approaches



# From standard to meta system identification

- System Identification
  - **Dataset:** I/O trajectory  $D = (u_{1:N}, y_{1:N})$  of one **particular** system
  - **Objective:** estimate a model from  $D$

$$\theta = \text{Alg}(D)$$
$$\hat{y}_{1:N}^* = M_{\theta}(u_{1:N}^*)$$



- Meta System Identification
  - **Meta dataset:** collection of datasets  $\mathcal{D} = \{D^{(1)}, D^{(2)}, \dots\}$  from **similar** systems
  - **Objective:** learn to make prediction on all systems in  $\mathcal{D}$

model-based meta learning


$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$



in-context learning


$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$

- Alg **is learned!**



# In-context system identification

- Particular dynamical system:
  - Input/output **context** data, input **query**  $\rightarrow$  predicted query output:



# Large Language Model analogy

## Output

- *Milan: Santa Maria delle Grazie, where you can see Leonardo da Vinci's 'The Last Supper'. The Milan Cathedral (Duomo di Milano) is another must-visit for its breathtaking Gothic architecture.*
- *Turin: Royal Palace of Turin, which offers a glimpse into the lives of the Italian monarchy. The Egyptian Museum in Turin, is also a must for history buffs.*





## Context

*A user is passionate about history and is planning a trip to North Italy.*

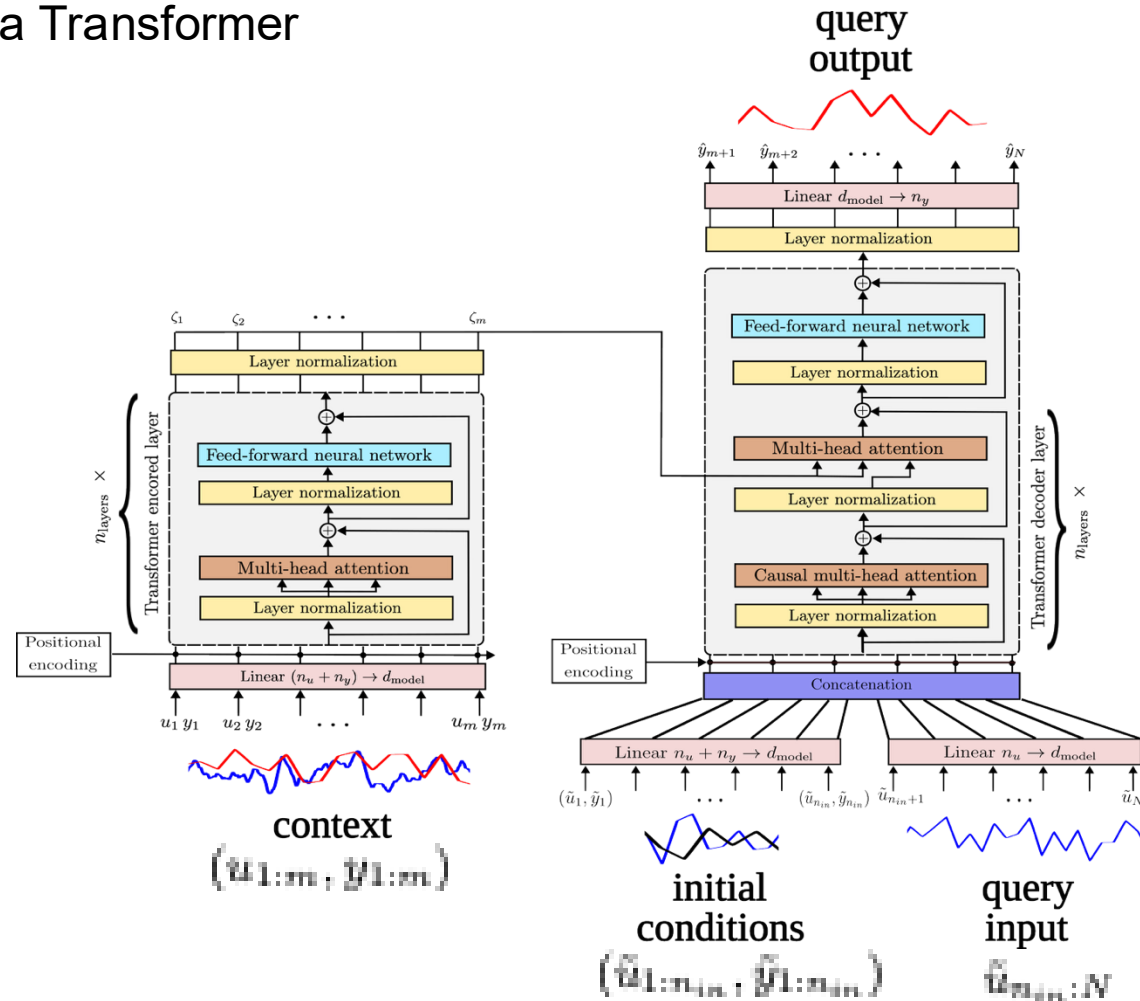
## Query

*Can you recommend some places to visit?*



# Meta model architecture

- Meta-model extended to process the query's input/output initial condition
- Parameterized as a Transformer





# Meta model training

## STEPS

1. Collect a meta dataset of similar systems from a **simulator**:

$$\mathfrak{D} = (D^{(1)}, D^{(2)}, \dots), \quad i=1, 2, \dots, \infty$$

2. Meta training **offline** to estimate a meta-model  $\mathcal{M}_\phi$ :


$$\hat{\phi} = \arg \min_{\phi} \sum_i \text{MSE}(D^{(i)}, \mathcal{M}_\phi)$$

3. Make predictions on new (real) system, in a **zero-shot** fashion (no training):

$$\hat{\mathbf{y}}^* = \mathcal{M}_{\hat{\phi}}(\mathbf{u}^*, \mathbf{D})$$

- Technically, it is just a supervised learning problem – standard training.
- The trained  $\mathcal{M}$  may be seen as a **learned** identification + simulation algorithm...
- ... or as a *specialized* question answering system – a kind of SYSID GPT.

### Wiener-Hammerstein Systems


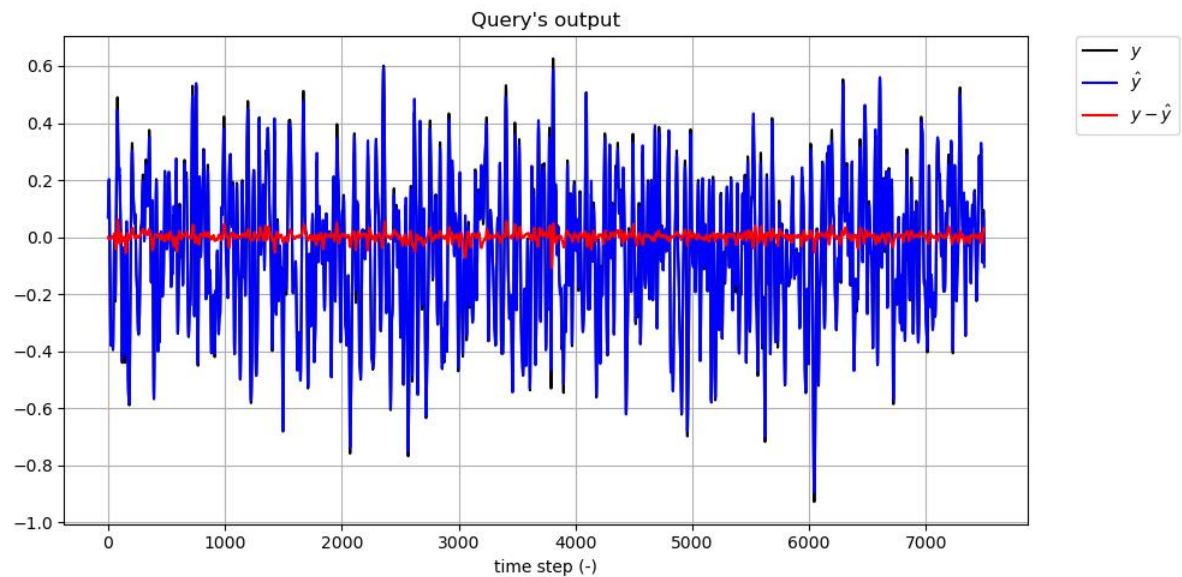


- Sequential LTI  $\rightarrow F(\cdot) \rightarrow$  LTI
- Random LTI, order  $\leq 5$
- $F(\cdot)$ : random feed-forward NN

- Meta training on **~100 M simulated** Wiener-Hammerstein systems
- The input  $u_{1:N}$  is a multisine signal with random phase and spectrum
- White noise added to the simulated output
- Transformer model with **~5 M parameters**, training time **~1 day** on an Nvidia 4090 GPU...

# Testing – Nonlinear benchmarks (www.nonlinearbenchmark.org)

- Real Wiener-Hammerstein benchmark




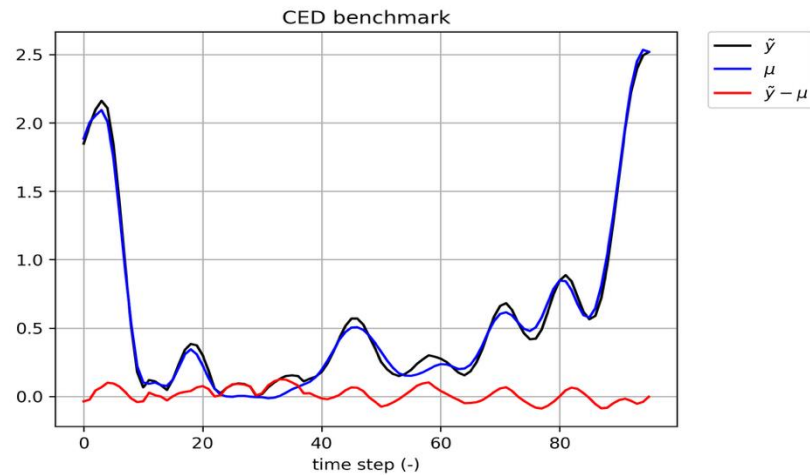
RMSE * 1000	Inference time (s)
1.4	5.91

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgiione, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

## Testing – Nonlinear benchmarks

- Coupled Electric Drives (CED)


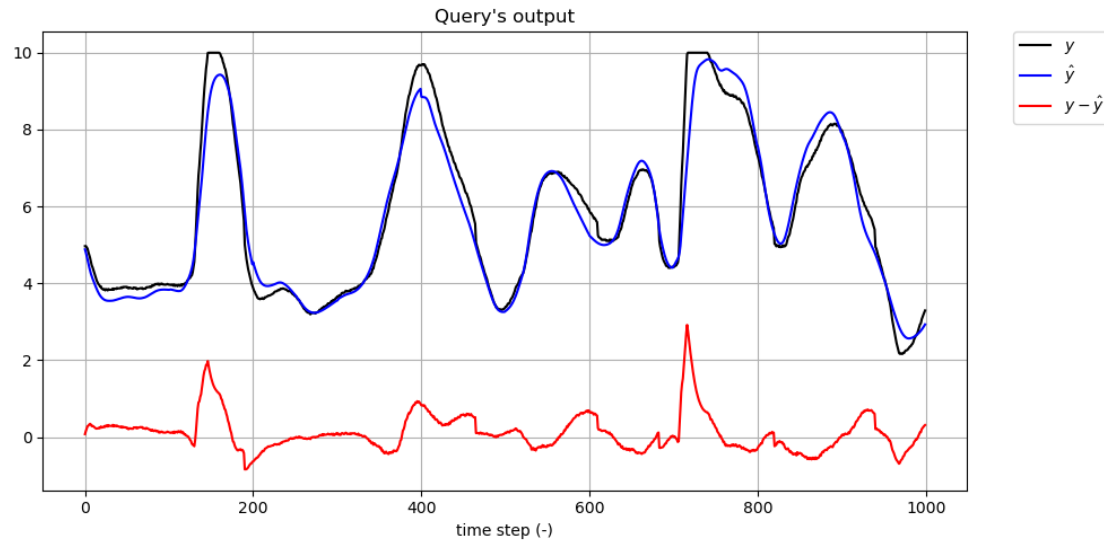


RMSE	Inference time (s)
[0.068,0.055]	0.17

State-of-the-art

## Testing – Nonlinear benchmarks

- Cascaded tanks with Overflow





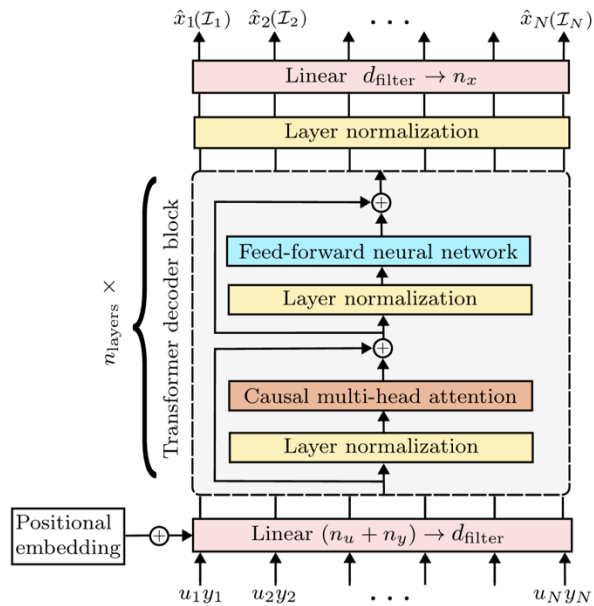
RMSE	Inference time (s)
0.429	0.22

State-of-the-art

Rufolo, M., Piga, D., Pura, F., & Forgone, M. (2025). In-Context Learning for System Identification: Recent Advances. Submitted to a journal, available [here](#).

## Other objectives: meta state estimation

- A meta state estimator that generalizes to different systems from the context



- Experiments on speed estimation of BLDC motors

Colombo, A., Busetto, R., Breschi, V., Forgone, M., Piga, D., & Formentin, S. (2025). In-Context Learning for Zero-Shot Speed Estimation of BLDC motors. arXiv preprint arXiv:2504.00673.

Busetto, R., Breschi, V., Forgone, M., Piga, D., & Formentin, S. (2024). In-context learning of state estimators. IFAC-PapersOnLine, 58(15), 145-150. Presented at SYSID 2024.

## Whitening the meta-learning box

The meta-learning grayscale

model-based meta learning

$$\hat{y}_{1:N}^* = M_{\text{Alg}_{\varphi}(D)}(u_{1:N}^*)$$



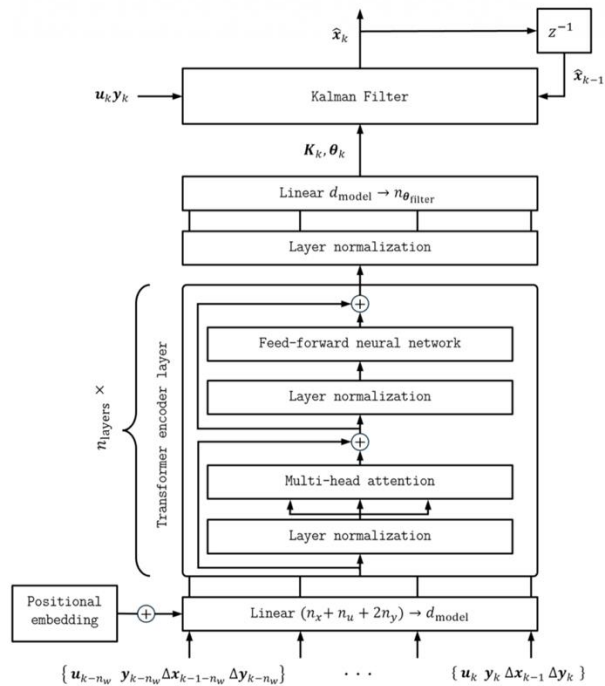
in-context learning

$$\hat{y}_{1:N}^* = \mathcal{M}_{\varphi}(u_{1:N}^*, D)$$

- In-context learning: prediction is a **black-box function** of dataset and new data.
- Other extreme: **learn hyper-parameters**  $\varphi$  of algorithm  $\text{Alg}_{\varphi}$  that returns (interpretable) **model parameters**
  - Learn the learning rate, initial value for gradient-based optimization (MAML), etc

## Meta state estimation: whitening the box

- Transformer returns estimator gain  $K$  + **physical** parameters  $\theta$  instead of predictions directly.



Model-based meta learning



In-context learning

- On-going work with Université de Poitiers (Hugo Koide, Guillaume Mercère)



# Conclusions

## Takeaways:

- In-context system identification, estimation, & control
- Off-line: learn how to map data into prediction
- On-line: provide context data, query input → query output ([zero shot](#))

## Current activities:

- Real-world applications
- Whitening the meta-learning box
- Providing guarantees